

#Predict the price of the Uber ride from a given pickup point to the agreed drop-off location. Perform following tasks:

1. Pre-process the dataset.
2. Identify outliers.
3. Check the correlation.
4. Implement linear regression and random forest regression models.
5. Evaluate the models and compare their respective scores like R2, RMSE, etc. Dataset link: <https://www.kaggle.com/datasets/yasserh/uber-fares-dataset>

```
In [1]: #Importing the required libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: #importing the dataset
df = pd.read_csv("uber.csv")
```

1. Pre-process the dataset.

```
In [3]: df.head()
```

```
Out[3]:
```

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354	-73.999512	40.723217	1
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	-73.994710	40.750325	1
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770	-73.962565	40.772647	1
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	-73.965316	40.803349	3
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085	-73.973082	40.761247	5

```
In [4]: df.info() #To get the required information of the dataset
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
Unnamed: 0      200000 non-null int64
key             200000 non-null object
fare_amount     200000 non-null float64
pickup_datetime 200000 non-null object
pickup_longitude 200000 non-null float64
pickup_latitude 200000 non-null float64
dropoff_longitude 199999 non-null float64
dropoff_latitude 199999 non-null float64
passenger_count 200000 non-null int64
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
```

```
In [5]: df.columns #TO get number of columns in the dataset
```

```
Out[5]: Index(['Unnamed: 0', 'key', 'fare_amount', 'pickup_datetime',
              'pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
              'dropoff_latitude', 'passenger_count'],
              dtype='object')
```

```
In [6]: df = df.drop(['Unnamed: 0', 'key'], axis= 1) #To drop unnamed column as it isn't required
```

```
In [7]: df.head()
```

```
Out[7]:
```

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
0	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354	-73.999512	40.723217	1
1	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	-73.994710	40.750325	1
2	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770	-73.962565	40.772647	1
3	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	-73.965316	40.803349	3
4	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085	-73.973082	40.761247	5

```
In [8]: df.shape #To get the total (Rows,Columns)
```

```
Out[8]: (200000, 7)
```

```
In [9]: df.dtypes #To get the type of each column
```

```
Out[9]: fare_amount      float64
pickup_datetime    object
pickup_longitude   float64
pickup_latitude    float64
dropoff_longitude  float64
dropoff_latitude   float64
passenger_count    int64
```

Column pickup_datetime is in wrong format (Object). Convert it to DateTime Format

```
In [10]: df.pickup_datetime = pd.to_datetime(df.pickup_datetime)
```

```
In [11]: df.dtypes
```

```
Out[11]: fare_amount          float64
pickup_datetime    datetime64[ns, UTC]
pickup_longitude    float64
pickup_latitude     float64
dropoff_longitude    float64
dropoff_latitude    float64
passenger_count     int64
dtype: object
```

Filling Missing values

```
In [12]: df.isnull().sum()
```

```
Out[12]: fare_amount          0
pickup_datetime             0
pickup_longitude            0
pickup_latitude             0
dropoff_longitude           1
dropoff_latitude            1
passenger_count             0
dtype: int64
```

```
In [13]: df['dropoff_latitude'].fillna(value=df['dropoff_latitude'].mean(),inplace = True)
df['dropoff_longitude'].fillna(value=df['dropoff_longitude'].median(),inplace = True)
```

```
In [14]: df.isnull().sum()
```

```
Out[14]: fare_amount          0
pickup_datetime             0
pickup_longitude            0
pickup_latitude             0
dropoff_longitude           0
dropoff_latitude            0
passenger_count             0
dtype: int64
```

To segregate each time of date and time

```
In [15]: df= df.assign(hour = df.pickup_datetime.dt.hour,
                        day= df.pickup_datetime.dt.day,
                        month = df.pickup_datetime.dt.month,
                        year = df.pickup_datetime.dt.year,
                        dayofweek = df.pickup_datetime.dt.dayofweek)
```

```
In [16]: df.head()
```

```
Out[16]:
```

	fare_amount		pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	hour	day	month	year	dayofweek
0	7.5	2015-05-07 19:52:06+00:00	-73.999817	40.738354	-73.999512	40.723217	1	19	7	5	2015	3	
1	7.7	2009-07-17 20:04:56+00:00	-73.994355	40.728225	-73.994710	40.750325	1	20	17	7	2009	4	
2	12.9	2009-08-24 21:45:00+00:00	-74.005043	40.740770	-73.962565	40.772647	1	21	24	8	2009	0	
3	5.3	2009-06-26 08:22:21+00:00	-73.976124	40.790844	-73.965316	40.803349	3	8	26	6	2009	4	
4	16.0	2014-08-28 17:47:00+00:00	-73.925023	40.744085	-73.973082	40.761247	5	17	28	8	2014	3	

Here we are going to use Haversine formula to calculate the distance between two points and journey, using longitude and latitude values.

Haversine formula $\text{hav}(\theta) = \sin^2(\theta/2)$.

```
In [19]: from math import *
# function to calculate the travel distance from the longitudes and latitudes
def distance_transform(longitude1, latitude1, longitude2, latitude2):
    travel_dist = []

    for pos in range(len(longitude1)):
        long1,lat1, long2, lat2 = map(radians,[longitude1[pos],latitude1[pos],longitude2[pos],latitude2[pos]])
        dist_long = long2 - long1
        dist_lat1 = lat2 - lat1
        a = sin(dist_lat1/2)**2 + cos(lat1) * cos(lat2) * sin(dist_long/2)**2
        c = 2 * asin(sqrt(a))*6371
        travel_dist.append(c)

    return travel_dist
```

```
in [21]: df.head()
```

```
out[21]:
```

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	hour	day	month	year	dayofweek	dist_travel_km
0	7.5	2015-05-07 19:52:06+00:00	-73.999817	40.738354	-73.999512	40.723217	1	19	7	5	2015	3	1.683323
1	7.7	2009-07-17 20:04:56+00:00	-73.994355	40.728225	-73.994710	40.750325	1	20	17	7	2009	4	2.457590
2	12.9	2009-08-24 21:45:00+00:00	-74.005043	40.740770	-73.962565	40.772647	1	21	24	8	2009	0	5.036377
3	5.3	2009-06-26 08:22:21+00:00	-73.976124	40.790844	-73.965316	40.803349	3	8	26	6	2009	4	1.661683
4	16.0	2014-08-28 17:47:00+00:00	-73.925023	40.744085	-73.973082	40.761247	5	17	28	8	2014	3	4.475450

```
in [22]: # drop the column 'pickup_datetime' using drop()
# 'axis = 1' drops the specified column

df = df.drop('pickup_datetime',axis=1)
```

```
in [23]: df.head()
```

```
out[23]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	hour	day	month	year	dayofweek	dist_travel_km
0	7.5	-73.999817	40.738354	-73.999512	40.723217	1	19	7	5	2015	3	1.683323
1	7.7	-73.994355	40.728225	-73.994710	40.750325	1	20	17	7	2009	4	2.457590
2	12.9	-74.005043	40.740770	-73.962565	40.772647	1	21	24	8	2009	0	5.036377
3	5.3	-73.976124	40.790844	-73.965316	40.803349	3	8	26	6	2009	4	1.661683
4	16.0	-73.925023	40.744085	-73.973082	40.761247	5	17	28	8	2014	3	4.475450


Checking outliers and filling them


```
in [24]: df.plot(kind = "box",subplots = True,layout = (7,2),figsize=(15,20)) #Boxplot to check the outliers
```

```
out[24]:
```

fare_amount	AxesSubplot(0.125,0.787927;0.352273x0.0920732)
pickup_longitude	AxesSubplot(0.547727,0.787927;0.352273x0.0920732)
pickup_latitude	AxesSubplot(0.125,0.677439;0.352273x0.0920732)
dropoff_longitude	AxesSubplot(0.547727,0.677439;0.352273x0.0920732)
dropoff_latitude	AxesSubplot(0.125,0.566951;0.352273x0.0920732)
passenger_count	AxesSubplot(0.547727,0.566951;0.352273x0.0920732)
hour	AxesSubplot(0.125,0.456463;0.352273x0.0920732)
day	AxesSubplot(0.547727,0.456463;0.352273x0.0920732)
month	AxesSubplot(0.125,0.345976;0.352273x0.0920732)
year	AxesSubplot(0.547727,0.345976;0.352273x0.0920732)
dayofweek	AxesSubplot(0.125,0.235488;0.352273x0.0920732)
dist_travel_km	AxesSubplot(0.547727,0.235488;0.352273x0.0920732)

dtype: object






```
In [25]: #Using the InterQuartile Range to fill the values
def remove_outlier(df1 , col):
    Q1 = df1[col].quantile(0.25)
    Q3 = df1[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_whisker = Q1-1.5*IQR
    upper_whisker = Q3+1.5*IQR
    df[col] = np.clip(df1[col] , lower_whisker , upper_whisker)
    return df1


def treat_outliers_all(df1 , col_list):
    for c in col_list:
        df1 = remove_outlier(df , c)
    return df1
```

```
In [26]: df = treat_outliers_all(df , df.iloc[:, 0::])
```

```
In [27]: df.plot(kind = "box",subplots = True,layout = (7,2),figsize=(15,20)) #Boxplot shows that dataset is free from outliers
```

```
Out[27]: fare_amount      AxesSubplot(0.125,0.787927;0.352273x0.0920732)
pickup_longitude    AxesSubplot(0.547727,0.787927;0.352273x0.0920732)
pickup_latitude      AxesSubplot(0.125,0.677439;0.352273x0.0920732)
dropoff_longitude    AxesSubplot(0.547727,0.677439;0.352273x0.0920732)
dropoff_latitude     AxesSubplot(0.125,0.566951;0.352273x0.0920732)
passenger_count     AxesSubplot(0.547727,0.566951;0.352273x0.0920732)
hour                AxesSubplot(0.125,0.456463;0.352273x0.0920732)
day                AxesSubplot(0.547727,0.456463;0.352273x0.0920732)
month              AxesSubplot(0.125,0.345976;0.352273x0.0920732)
year              AxesSubplot(0.547727,0.345976;0.352273x0.0920732)
dayofweek         AxesSubplot(0.125,0.235488;0.352273x0.0920732)
dist_travel_km    AxesSubplot(0.547727,0.235488;0.352273x0.0920732)
dtype: object
```





```
In [28]: #Uber doesn't travel over 130 kms so minimize the distance
df= df.loc[(df.dist_travel_km >= 1) | (df.dist_travel_km <= 130)]
print("Remaining observations in the dataset:", df.shape)

Remaining observations in the dataset: (200000, 12)
```

```
In [29]: #Finding incorrect latitude (less than or greater than 90) and longitude (greater than or less than 180)
incorrect_coordinates = df.loc[(df.pickup_latitude > 90) | (df.pickup_latitude < -90) |
(df.dropoff_latitude > 90) | (df.dropoff_latitude < -90) |
(df.pickup_longitude > 180) | (df.pickup_longitude < -180) |
(df.dropoff_longitude > 90) | (df.dropoff_longitude < -90)
]
```

```
In [30]: df.drop(incorrect_coordinates, inplace = True, errors = 'ignore')
```

```
In [31]: df.head()
```

```
Out[31]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	hour	day	month	year	dayofweek	dist_travel_km
0	7.5	-73.999817	40.738354	-73.999512	40.723217	1.0	19	7	5	2015	3	1.683323
1	7.7	-73.994355	40.728225	-73.994710	40.750325	1.0	20	17	7	2009	4	2.457590
2	12.9	-74.005043	40.740770	-73.962565	40.772647	1.0	21	24	8	2009	0	5.036377
3	5.3	-73.976124	40.790844	-73.965316	40.803349	3.0	8	26	6	2009	4	1.661683
4	16.0	-73.929786	40.744085	-73.973082	40.761247	3.5	17	28	8	2014	3	4.475450


```
In [32]: df.isnull().sum()
```

```
Out[32]:
```

fare_amount	0
pickup_longitude	0
pickup_latitude	0
dropoff_longitude	0
dropoff_latitude	0
passenger_count	0
hour	0
day	0
month	0
year	0
dayofweek	0
dist_travel_km	0
dtype:	int64

```
In [33]: sns.heatmap(df.isnull()) #Free for null values
```

```
Out[33]: <matplotlib.axes._subplots.AxesSubplot at 0x8d8af2a080>
```



```
In [34]: corr = df.corr() #Function to find the correlation
```

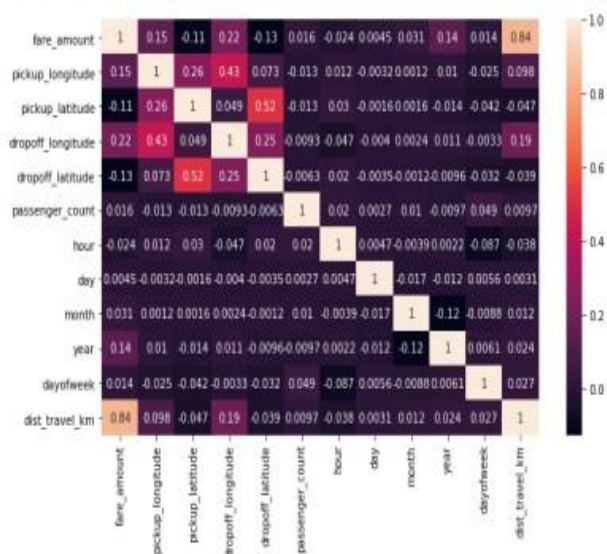
```
In [35]: corr
```

```
Out[35]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	hour	day	month	year	dayofweek	dist_travel_km
fare_amount	1.000000	0.154069	-0.110842	0.218675	-0.125898	0.015778	-0.023623	0.004534	0.030817	0.141277	0.013652	0.844374
pickup_longitude	0.154069	1.000000	0.259497	0.425619	0.073290	-0.013213	0.011579	-0.003204	0.001169	0.010198	-0.024652	0.098094
pickup_latitude	-0.110842	0.259497	1.000000	0.048889	0.515714	-0.012889	0.029681	-0.001553	0.001562	-0.014243	-0.042310	-0.046812
dropoff_longitude	0.218675	0.425619	0.048889	1.000000	0.245667	-0.009303	-0.046558	-0.004007	0.002391	0.011346	-0.003336	0.186531
dropoff_latitude	-0.125898	0.073290	0.515714	0.245667	1.000000	-0.006308	0.019783	-0.003479	-0.001193	-0.009603	-0.031919	-0.038900
passenger_count	0.015778	-0.013213	-0.012889	-0.009303	-0.006308	1.000000	0.020274	0.002712	0.010351	-0.009749	0.048550	0.009709
hour	-0.023623	0.011579	0.029681	-0.046558	0.019783	0.020274	1.000000	0.004677	-0.003926	0.002156	-0.086947	-0.038366
day	0.004534	-0.003204	-0.001553	-0.004007	-0.003479	0.002712	0.004677	1.000000	-0.017360	-0.012170	0.005617	0.003062
month	0.030817	0.001169	0.001562	0.002391	-0.001193	0.010351	-0.003926	-0.017360	1.000000	-0.115859	-0.008786	0.011628
year	0.141277	0.010198	-0.014243	0.011346	-0.009603	-0.009749	0.002156	-0.012170	-0.115859	1.000000	0.006113	0.024278
dayofweek	0.013652	-0.024652	-0.042310	-0.003336	-0.031919	0.048550	-0.086947	0.005617	-0.008786	0.006113	1.000000	0.027053
dist_travel_km	0.844374	0.098094	-0.046812	0.186531	-0.038900	0.009709	-0.038366	0.003062	0.011628	0.024278	0.027053	1.000000

```
In [36]: fig,axis = plt.subplots(figsize = (10,6))
sns.heatmap(df.corr(),annot = True) #Correlation Heatmap (Light values means highly correlated)
```

```
Out[36]: <matplotlib.axes._subplots.AxesSubplot at 0x8d8affc588>
```



Dividing the dataset into feature and target values

```
In [37]: x = df[['pickup_longitude','pickup_latitude','dropoff_longitude','dropoff_latitude','passenger_count','hour','day','month','year','dayofweek','dist_travel_km']]
```

```
In [38]: y = df['fare_amount']
```

Dividing the dataset into training and testing dataset

```
In [39]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(x,y,test_size = 0.33)
```

Linear Regression

```
In [40]: from sklearn.linear_model import LinearRegression
regression = LinearRegression()
```

```
In [41]: regression.fit(X_train,y_train)
```

```
Out[41]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
In [42]: regression.intercept_ #To find the linear intercept
```

```
Out[42]: 2809.192377415925
```

```
In [43]: regression.coef_ #To find the linear coefficient
```

```
In [44]: prediction = regression.predict(X_test) #To predict the target values
```

```
In [45]: print(prediction)
```

```
[18.88422892  4.74787896  9.95283165 ...  5.89597937 17.88144322
 5.38487972]
```

```
In [46]: y_test
```

```
Out[46]: 16858      8.50
181876     4.10
78798      9.30
87421     12.90
169443     22.25
18976     11.00
58921     13.70
199564     14.50
125215     5.30
67510      8.50
85217     22.25
156983     21.50
116795     4.10
112179     16.90
124459     3.70
173299     22.25
51448     19.70
99582     22.25
174467     18.90
78880     20.50
26798     22.25
38581     4.50
63091     12.90
171287     22.25
142238     8.50
101186     7.30
128177     4.50
154585     14.50
75840      5.50
85918     14.00
...
184227     18.10
14172     19.70
49985      3.70
183845     6.50
11927     12.90
93684      4.50
181795     13.70
21444      6.10
85147      8.50
81311      8.00
157686     11.70
194874      6.50
132558     18.50
132616     11.70
188536      5.70
179629      8.90
11277      3.70
147880      7.30
116553      5.70
157394      6.50
183519     13.30
41348     12.90
12688      4.50
6828       5.50
84612      5.00
168836      3.70
39719     21.00
124536      4.90
98432     22.10
12543      4.90
```

```
Name: fare_amount, Length: 66000, dtype: float64
```

```

147880    7.30
116553    5.70
157394    6.50
103519   13.30
41348    12.90
12608     4.50
6820     5.50
84612     5.00
168836    3.70
39719    21.00
124536    4.90
90432    22.10
12543     4.90
Name: fare_amount, Length: 66000, dtype: float64

```

Metrics Evaluation using R2, Mean Squared Error, Root Mean Squared Error

```

In [47]: from sklearn.metrics import r2_score

In [48]: r2_score(y_test, prediction)

Out[48]: 0.7471032194200018

In [49]: from sklearn.metrics import mean_squared_error

In [50]: MSE = mean_squared_error(y_test, prediction)

In [51]: MSE

Out[51]: 7.464818887848474

In [52]: RMSE = np.sqrt(MSE)

In [53]: RMSE

Out[53]: 2.7321820744321696

```

Random Forest Regression

```

In [54]: from sklearn.ensemble import RandomForestRegressor

In [55]: rf = RandomForestRegressor(n_estimators=100) #Here n_estimators means number of trees you want to build before making the prediction

In [56]: rf.fit(X_train, y_train)

Out[56]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                               max_features='auto', max_leaf_nodes=None,
                               min_impurity_decrease=0.0, min_impurity_split=None,
                               min_samples_leaf=1, min_samples_split=2,
                               min_weight_fraction_leaf=0.0, n_estimators=100,
                               n_jobs=None, oob_score=False, random_state=None,
                               verbose=0, warm_start=False)

In [57]: y_pred = rf.predict(X_test)

In [58]: y_pred

Out[58]: array([ 9.7025,  4.744 ,  9.202 , ...,  6.408 , 16.2802,  4.47  ])

```

Metrics evaluatin for Random Forest

```

In [59]: R2_Random = r2_score(y_test, y_pred)

In [60]: R2_Random

Out[60]: 0.8024361566950065

In [64]: MSE_Random = mean_squared_error(y_test, y_pred)
MSE_Random

Out[64]: 5.831542440662031

In [65]: RMSE_Random = np.sqrt(MSE_Random)
RMSE_Random

Out[65]: 2.4148586792319815

```


Assignment 2

1. Classify the email using the binary classification method. Email Spam detection has two states:

- Normal State – Not Spam,
- Abnormal State – Spam. Use K-Nearest Neighbors and Support Vector Machine for classification. Analyze their performance.

Dataset link: The emails.csv dataset on the Kaggle <https://www.kaggle.com/datasets/balaka18/email-spam-classification-dataset-csv>

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn import metrics
```

```
In [9]: df=pd.read_csv("C:\\Users\\Shree\\Downloads\\archive\\emails.csv")
```

```
In [11]: df.head()
```

```
Out[11]:
```

	Email No.	the	to	ect	and	for	of	a	you	hou	_	convey	jay	valued	lay	infrastructure	military	allowing	ff	dry	Prediction
0	Email 1	0	0	1	0	0	0	2	0	0	_	0	0	0	0	0	0	0	0	0	0
1	Email 2	8	13	24	6	6	2	102	1	27	_	0	0	0	0	0	0	0	1	0	0
2	Email 3	0	0	1	0	0	0	8	0	0	_	0	0	0	0	0	0	0	0	0	0
3	Email 4	0	5	22	0	5	1	51	2	10	_	0	0	0	0	0	0	0	0	0	0
4	Email 5	7	6	17	1	5	2	57	0	9	_	0	0	0	0	0	0	0	1	0	0

5 rows × 3002 columns

```
In [12]: df.columns
```

```
Out[12]: Index(['Email No.', 'the', 'to', 'ect', 'and', 'for', 'of', 'a', 'you', 'hou',
...
'convey', 'jay', 'valued', 'lay', 'infrastructure', 'military',
'allowing', 'ff', 'dry', 'Prediction'],
dtype='object', length=3002)
```

```
In [13]: df.isnull().sum()
```

```
Out[13]: Email No.      0
the              0
to              0
ect             0
and             0
..
military        0
allowing        0
ff              0
dry             0
Prediction      0
Length: 3002, dtype: int64
```

```
In [14]: df.isnull().sum()
```

```
Out[14]: Email No.      0
the              0
to              0
ect             0
and             0
..
military        0
allowing        0
ff              0
dry             0
Prediction      0
Length: 3002, dtype: int64
```

```
In [15]: df.drop(['Email No.'],axis=1,inplace=True)
X = df.drop(['Prediction'],axis = 1)
```

In []: KNN classifier

```
In [16]: from sklearn.preprocessing import scale
X = scale(X)
# split into train and test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 42)
```

```
In [17]: from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=7)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
```

```
In [18]: print("Prediction",y_pred)
```

Prediction [0 0 1 ... 1 1 1]

```
In [20]: print("KNN accuracy = ",metrics.accuracy_score(y_test,y_pred))
```

KNN accuracy = 0.8009020618556701

```
In [21]: print("Confusion matrix",metrics.confusion_matrix(y_test,y_pred))
```

Confusion matrix [[804 293]
[16 439]]

In []: SVM classifier

```
In [22]: # cost C = 1
model = SVC(C = 1)
# fit
model.fit(X_train, y_train)
# predict
y_pred = model.predict(X_test)
```

```
In [23]: metrics.confusion_matrix(y_true=y_test, y_pred=y_pred)
```

```
Out[23]: array([[1091,    6],
               [  90,  365]], dtype=int64)
```

```
In [24]: print("SVM accuracy = ",metrics.accuracy_score(y_test,y_pred))
```

SVM accuracy = 0.9381443298969072

Dataset Description: The case study is from an open-source dataset from Kaggle. The dataset contains 10,000 sample points with 14 distinct features such as CustomerId, CreditScore, Geography, Gender, Age, Tenure, Balance, etc. Link to the Kaggle project: <https://www.kaggle.com/barelydedicated/bank-customer-churn-modeling> Perform following steps:

- ```
In [46]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt #Importing the libraries
```

[illegible]

```
In [51]: df.isnull()
```

```
Out[51]:
```

|      | RowNumber | CustomerId | Surname | CreditScore | Geography | Gender | Age   | Tenure | Balance | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSalary |
|------|-----------|------------|---------|-------------|-----------|--------|-------|--------|---------|---------------|-----------|----------------|-----------------|
| 0    | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False           |
| 1    | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False           |
| 2    | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False           |
| 3    | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False           |
| 4    | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False           |
| ...  | ...       | ...        | ...     | ...         | ...       | ...    | ...   | ...    | ...     | ...           | ...       | ...            | ...             |
| 9995 | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False           |
| 9996 | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False           |
| 9997 | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False           |
| 9998 | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False           |
| 9999 | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False           |

10000 rows x 14 columns

```
In [52]: df.isnull().sum()
```

```
Out[52]: RowNumber 0
CustomerId 0
Surname 0
CreditScore 0
Geography 0
Gender 0
Age 0
Tenure 0
Balance 0
NumOfProducts 0
HasCrCard 0
IsActiveMember 0
EstimatedSalary 0
Exited 0
dtype: int64
```

```
In [53]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
Column Non-Null Count Dtype
--- ---
0 RowNumber 10000 non-null int64
1 CustomerId 10000 non-null int64
2 Surname 10000 non-null object
3 CreditScore 10000 non-null int64
4 Geography 10000 non-null object
5 Gender 10000 non-null object
6 Age 10000 non-null int64
7 Tenure 10000 non-null int64
8 Balance 10000 non-null float64
9 NumOfProducts 10000 non-null int64
10 HasCrCard 10000 non-null int64
11 IsActiveMember 10000 non-null int64
12 EstimatedSalary 10000 non-null float64
13 Exited 10000 non-null int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
In [54]: df.dtypes
```

```
Out[54]: RowNumber int64
CustomerId int64
Surname object
CreditScore int64
Geography object
Gender object
Age int64
Tenure int64
Balance float64
NumOfProducts int64
HasCrCard int64
IsActiveMember int64
EstimatedSalary float64
Exited int64
dtype: object
```

```
In [55]: df.columns
```



```
In [55]: df.columns
```

```
Out[55]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
 'IsActiveMember', 'EstimatedSalary', 'Exited'],
 dtype='object')
```

```
In [56]: df = df.drop(['RowNumber', 'Surname', 'CustomerId'], axis= 1) #Dropping the unnecessary columns
```

```
In [57]: df.head()
```

```
Out[57]:
```


|   | CreditScore | Geography | Gender | Age | Tenure | Balance   | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSalary | Exited |
|---|-------------|-----------|--------|-----|--------|-----------|---------------|-----------|----------------|-----------------|--------|
| 0 | 619         | France    | Female | 42  | 2      | 0.00      | 1             | 1         | 1              | 101348.88       | 1      |
| 1 | 608         | Spain     | Female | 41  | 1      | 83907.86  | 1             | 0         | 1              | 112542.58       | 0      |
| 2 | 502         | France    | Female | 42  | 8      | 159660.80 | 3             | 1         | 0              | 113931.57       | 1      |
| 3 | 699         | France    | Female | 39  | 1      | 0.00      | 2             | 0         | 0              | 93826.63        | 0      |
| 4 | 850         | Spain     | Female | 43  | 2      | 125510.82 | 1             | 1         | 1              | 79084.10        | 0      |

## Visualization

```
In [101]: def visualization(x, y, xlabel):
 plt.figure(figsize=(10,5))
 plt.hist([x, y], color=['red', 'green'], label = ['exit', 'not_exit'])
 plt.xlabel(xlabel,fontsize=20)
 plt.ylabel("No. of customers", fontsize=20)
 plt.legend()
```


```
In [102]: df_churn_exited = df[df['Exited']==1]['Tenure']
df_churn_not_exited = df[df['Exited']==0]['Tenure']
```

```
In [103]: visualization(df_churn_exited, df_churn_not_exited, "Tenure")
```



```
In [105]: df_churn_exited2 = df[df['Exited']==1]['Age']
df_churn_not_exited2 = df[df['Exited']==0]['Age']
```

```
In [106]: visualization(df_churn_exited2, df_churn_not_exited2, "Age")
```



```
In [75]: classifier.add(Dense(activation = "relu",units = 6,kernel_initializer = "uniform")) #Adding second hidden layers
```

```
In [76]: classifier.add(Dense(activation = "sigmoid",units = 1,kernel_initializer = "uniform")) #Final neuron will be having sigmoid function
```

```
In [77]: classifier.compile(optimizer="adam",loss = 'binary_crossentropy',metrics = ['accuracy']) #To compile the Artificial Neural Network. Used Binary crossentropy as we just have only two output
```

```
In [79]: classifier.summary() #3 layers created. 6 neurons in 1st,6neurons in 2nd layer and 1 neuron in last
```

Model: "sequential\_1"

| Layer (type)            | Output Shape | Param # |
|-------------------------|--------------|---------|
| dense_3 (Dense)         | (None, 6)    | 72      |
| dense_4 (Dense)         | (None, 6)    | 42      |
| dense_5 (Dense)         | (None, 1)    | 7       |
| Total params: 121       |              |         |
| Trainable params: 121   |              |         |
| Non-trainable params: 0 |              |         |

```
In [89]: classifier.fit(X_train,y_train,batch_size=10,epochs=50) #Fitting the ANN to training dataset
```

```
Epoch 1/50
700/700 [=====] - 0s 674us/step - loss: 0.4293 - accuracy: 0.7947
Epoch 2/50
700/700 [=====] - 0s 647us/step - loss: 0.4239 - accuracy: 0.7947
Epoch 3/50
700/700 [=====] - 0s 657us/step - loss: 0.4203 - accuracy: 0.8067
Epoch 4/50
700/700 [=====] - 0s 664us/step - loss: 0.4167 - accuracy: 0.8260
Epoch 5/50
700/700 [=====] - 0s 674us/step - loss: 0.4153 - accuracy: 0.8287
Epoch 6/50
700/700 [=====] - 0s 653us/step - loss: 0.4137 - accuracy: 0.8310
Epoch 7/50
700/700 [=====] - 0s 658us/step - loss: 0.4125 - accuracy: 0.8317
Epoch 8/50
700/700 [=====] - 1s 842us/step - loss: 0.4116 - accuracy: 0.8306
Epoch 9/50
700/700 [=====] - 0s 671us/step - loss: 0.4103 - accuracy: 0.8331
Epoch 10/50
700/700 [=====] - 0s 682us/step - loss: 0.4100 - accuracy: 0.8326
Epoch 11/50
700/700 [=====] - 0s 690us/step - loss: 0.4093 - accuracy: 0.8337
Epoch 12/50
700/700 [=====] - 0s 688us/step - loss: 0.4087 - accuracy: 0.8339
Epoch 13/50
700/700 [=====] - 0s 675us/step - loss: 0.4081 - accuracy: 0.8341
Epoch 14/50
700/700 [=====] - 1s 722us/step - loss: 0.4071 - accuracy: 0.8331
Epoch 15/50
700/700 [=====] - 1s 811us/step - loss: 0.4065 - accuracy: 0.8341
Epoch 16/50
700/700 [=====] - 0s 711us/step - loss: 0.4056 - accuracy: 0.8356
Epoch 17/50
700/700 [=====] - 0s 702us/step - loss: 0.4046 - accuracy: 0.8366
Epoch 18/50
700/700 [=====] - 0s 688us/step - loss: 0.4035 - accuracy: 0.8343
Epoch 19/50
700/700 [=====] - 1s 715us/step - loss: 0.4024 - accuracy: 0.8363
Epoch 20/50
700/700 [=====] - 0s 714us/step - loss: 0.4020 - accuracy: 0.8337
Epoch 21/50
700/700 [=====] - 0s 705us/step - loss: 0.4010 - accuracy: 0.8374
Epoch 22/50
700/700 [=====] - 1s 720us/step - loss: 0.4003 - accuracy: 0.8370
Epoch 23/50
700/700 [=====] - 0s 692us/step - loss: 0.3993 - accuracy: 0.8374
Epoch 24/50
700/700 [=====] - 0s 709us/step - loss: 0.3990 - accuracy: 0.8356
Epoch 25/50
700/700 [=====] - 1s 871us/step - loss: 0.3984 - accuracy: 0.8366
Epoch 26/50
700/700 [=====] - 1s 719us/step - loss: 0.3984 - accuracy: 0.8367
Epoch 27/50
700/700 [=====] - 1s 719us/step - loss: 0.3980 - accuracy: 0.8366
Epoch 28/50
700/700 [=====] - 0s 695us/step - loss: 0.3981 - accuracy: 0.8366
Epoch 29/50
700/700 [=====] - 0s 667us/step - loss: 0.3976 - accuracy: 0.8374
Epoch 30/50
700/700 [=====] - 0s 669us/step - loss: 0.3972 - accuracy: 0.8373
```

```

700/700 [=====] - 1s 771us/step - loss: 0.3960 - accuracy: 0.8370
Epoch 37/50
700/700 [=====] - 1s 1ms/step - loss: 0.3963 - accuracy: 0.8366
Epoch 38/50
700/700 [=====] - 1s 764us/step - loss: 0.3962 - accuracy: 0.8373
Epoch 39/50
700/700 [=====] - 1s 823us/step - loss: 0.3950 - accuracy: 0.8384
Epoch 40/50
700/700 [=====] - 1s 759us/step - loss: 0.3956 - accuracy: 0.8361
Epoch 41/50
700/700 [=====] - 1s 773us/step - loss: 0.3949 - accuracy: 0.8366
Epoch 42/50
700/700 [=====] - 0s 695us/step - loss: 0.3953 - accuracy: 0.8369
Epoch 43/50
700/700 [=====] - 0s 701us/step - loss: 0.3952 - accuracy: 0.8369
Epoch 44/50
700/700 [=====] - 0s 707us/step - loss: 0.3952 - accuracy: 0.8366
Epoch 45/50
700/700 [=====] - 0s 680us/step - loss: 0.3955 - accuracy: 0.8376
Epoch 46/50
700/700 [=====] - 0s 665us/step - loss: 0.3947 - accuracy: 0.8373
Epoch 47/50
700/700 [=====] - 0s 708us/step - loss: 0.3947 - accuracy: 0.8371
Epoch 48/50
700/700 [=====] - 0s 681us/step - loss: 0.3944 - accuracy: 0.8371
Epoch 49/50
700/700 [=====] - 0s 678us/step - loss: 0.3947 - accuracy: 0.8383
Epoch 50/50
700/700 [=====] - 1s 869us/step - loss: 0.3944 - accuracy: 0.8370

```

Out[89]: <tensorflow.python.keras.callbacks.History at 0x1fb1eb93df0>

```

In [90]: y_pred = classifier.predict(X_test)
 y_pred = (y_pred > 0.5) #Predicting the result

```

```

In [97]: from sklearn.metrics import confusion_matrix, accuracy_score, classification_report

```

```

In [92]: cm = confusion_matrix(y_test, y_pred)

```

```

In [93]: cm

```

Out[93]: array([[2328, 72],  
[ 425, 175]], dtype=int64)

```

In [94]: accuracy = accuracy_score(y_test, y_pred)

```

```

In [95]: accuracy

```


Out[95]: 0.8343333333333334

```

In [98]: plt.figure(figsize = (10,7))
 sns.heatmap(cm,annot = True)
 plt.xlabel('Predicted')
 plt.ylabel('Truth')

```

Out[98]: Text(69.0, 0.5, 'Truth')



```

In [100]: print(classification_report(y_test, y_pred))

```

```

precision recall f1-score support

```



# Implement K-Means clustering/ hierarchical clustering on sales\_data\_sample.csv dataset. Determine the number of clusters using the elbow method.

```
In [198]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
#Importing the required libraries.

In [199]: from sklearn.cluster import KMeans, k_means #For clustering
from sklearn.decomposition import PCA #Linear Dimensionality reduction.

In [200]: df = pd.read_csv("sales_data_sample.csv") #Loading the dataset.
```

## Preprocessing

```
In [201]: df.head()
```

```
Out[201]:
```

|   | ORDERNUMBER | QUANTITYORDERED | PRICEEACH | ORDERLINENUMBER | SALES   | ORDERDATE       | STATUS  | QTR_ID | MONTH_ID | YEAR_ID | ... | ADDRESSLINE1           |
|---|-------------|-----------------|-----------|-----------------|---------|-----------------|---------|--------|----------|---------|-----|------------------------|
| 0 | 10107       | 30              | 95.70     | 2               | 2871.00 | 2/24/2003 0:00  | Shipped | 1      | 2        | 2003    | ... | 897 Long Airpi Aven    |
| 1 | 10121       | 34              | 81.35     | 5               | 2765.90 | 5/7/2003 0:00   | Shipped | 2      | 5        | 2003    | ... | 59 rue l'Abba          |
| 2 | 10134       | 41              | 94.74     | 2               | 3884.34 | 7/1/2003 0:00   | Shipped | 3      | 7        | 2003    | ... | 27 rue Colonel Pier Au |
| 3 | 10145       | 45              | 83.26     | 6               | 3746.70 | 8/25/2003 0:00  | Shipped | 3      | 8        | 2003    | ... | 78934 Hillsi [         |
| 4 | 10159       | 49              | 100.00    | 14              | 5205.27 | 10/10/2003 0:00 | Shipped | 4      | 10       | 2003    | ... | 7734 Strong :          |

5 rows × 25 columns

```
In [202]: df.shape
```

```
Out[202]: (2823, 25)
```

```
In [203]: df.describe()
```

```
Out[203]:
```

|       | ORDERNUMBER  | QUANTITYORDERED | PRICEEACH   | ORDERLINENUMBER | SALES        | QTR_ID      | MONTH_ID    | YEAR_ID     | MSRP        |
|-------|--------------|-----------------|-------------|-----------------|--------------|-------------|-------------|-------------|-------------|
| count | 2823.000000  | 2823.000000     | 2823.000000 | 2823.000000     | 2823.000000  | 2823.000000 | 2823.000000 | 2823.000000 | 2823.000000 |
| mean  | 10258.725115 | 35.092809       | 83.658544   | 6.466171        | 3553.889072  | 2.717676    | 7.092455    | 2003.81509  | 100.715551  |
| std   | 92.085478    | 9.741443        | 20.174277   | 4.225841        | 1841.865106  | 1.203878    | 3.656633    | 0.69967     | 40.187912   |
| min   | 10100.000000 | 6.000000        | 26.880000   | 1.000000        | 482.130000   | 1.000000    | 1.000000    | 2003.00000  | 33.000000   |
| 25%   | 10180.000000 | 27.000000       | 68.860000   | 3.000000        | 2203.430000  | 2.000000    | 4.000000    | 2003.00000  | 68.000000   |
| 50%   | 10262.000000 | 35.000000       | 95.700000   | 6.000000        | 3184.800000  | 3.000000    | 8.000000    | 2004.00000  | 99.000000   |
| 75%   | 10333.500000 | 43.000000       | 100.000000  | 9.000000        | 4508.000000  | 4.000000    | 11.000000   | 2004.00000  | 124.000000  |
| max   | 10425.000000 | 97.000000       | 100.000000  | 18.000000       | 14082.800000 | 4.000000    | 12.000000   | 2005.00000  | 214.000000  |

```
In [204]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 25 columns):
Column Non-Null Count Dtype
--- -
0 ORDERNUMBER 2823 non-null int64
1 QUANTITYORDERED 2823 non-null int64
2 PRICEEACH 2823 non-null float64
```



```
In [205]: df.isnull().sum()
```

```
Out[205]: ORDERNUMBER 0
QUANTITYORDERED 0
PRICEEACH 0
ORDERLINENUMBER 0
SALES 0
ORDERDATE 0
STATUS 0
QTR_ID 0
MONTH_ID 0
YEAR_ID 0
PRODUCTLINE 0
MSRP 0
PRODUCTCODE 0
CUSTOMERNAME 0
PHONE 0
ADDRESSLINE1 0
ADDRESSLINE2 2521
CITY 0
STATE 1486
POSTALCODE 76
COUNTRY 0
TERRITORY 1074
CONTACTLASTNAME 0
CONTACTFIRSTNAME 0
DEALSIZE 0
dtype: int64
```

```
In [206]: df.dtypes
```

```
Out[206]: ORDERNUMBER int64
QUANTITYORDERED int64
PRICEEACH float64
ORDERLINENUMBER int64
SALES float64
ORDERDATE object
STATUS object
QTR_ID int64
MONTH_ID int64
YEAR_ID int64
PRODUCTLINE object
MSRP int64
PRODUCTCODE object
CUSTOMERNAME object
PHONE object
ADDRESSLINE1 object
ADDRESSLINE2 object
CITY object
STATE object
POSTALCODE object
COUNTRY object
TERRITORY object
CONTACTLASTNAME object
CONTACTFIRSTNAME object
DEALSIZE object
dtype: object
```

```
In [207]: df_drop = ['ADDRESSLINE1', 'ADDRESSLINE2', 'STATUS', 'POSTALCODE', 'CITY', 'TERRITORY', 'PHONE', 'STATE', 'CONTACTFI
RSTNAME', 'CONTACTLASTNAME', 'CUSTOMERNAME', 'ORDERNUMBER']
df = df.drop(df_drop, axis=1) #Dropping the categorical unnecessary columns along with columns having null values. C
an't fill the null values are there are alot of null values.
```

```
In [208]: df.isnull().sum()
```

```
Out[208]: QUANTITYORDERED 0
PRICEEACH 0
ORDERLINENUMBER 0
SALES 0
ORDERDATE 0
QTR_ID 0
MONTH_ID 0
```

```
In [209]: df.dtypes
```

```
Out[209]: QUANTITYORDERED int64
PRICEEACH float64
ORDERLINENUMBER int64
SALES float64
ORDERDATE object
QTR_ID int64
MONTH_ID int64
YEAR_ID int64
PRODUCTLINE object
MSRP int64
PRODUCTCODE object
COUNTRY object
DEALSIZE object
dtype: object
```

```
In []: # Checking the categorical columns.
```

```
In [210]: df['COUNTRY'].unique()
```

```
Out[210]: array(['USA', 'France', 'Norway', 'Australia', 'Finland', 'Austria', 'UK',
 'Spain', 'Sweden', 'Singapore', 'Canada', 'Japan', 'Italy',
 'Denmark', 'Belgium', 'Philippines', 'Germany', 'Switzerland',
 'Ireland'], dtype=object)
```

```
In [211]: df['PRODUCTLINE'].unique()
```

```
Out[211]: array(['Motorcycles', 'Classic Cars', 'Trucks and Buses', 'Vintage Cars',
 'Planes', 'Ships', 'Trains'], dtype=object)
```

```
In [212]: df['DEALSIZE'].unique()
```

```
Out[212]: array(['Small', 'Medium', 'Large'], dtype=object)
```

```
In [213]: productline = pd.get_dummies(df['PRODUCTLINE']) #Converting the categorical columns.
Dealsize = pd.get_dummies(df['DEALSIZE'])
```

```
In [214]: df = pd.concat([df, productline, Dealsize], axis = 1)
```

```
In [215]: df_drop = ['COUNTRY', 'PRODUCTLINE', 'DEALSIZE'] #Dropping Country too as there are alot of countries.
df = df.drop(df_drop, axis=1)
```

```
In [216]: df['PRODUCTCODE'] = pd.Categorical(df['PRODUCTCODE']).codes #Converting the datatype.
```

```
In [217]: df.drop('ORDERDATE', axis=1, inplace=True) #Dropping the Orderdate as Month is already included.
```

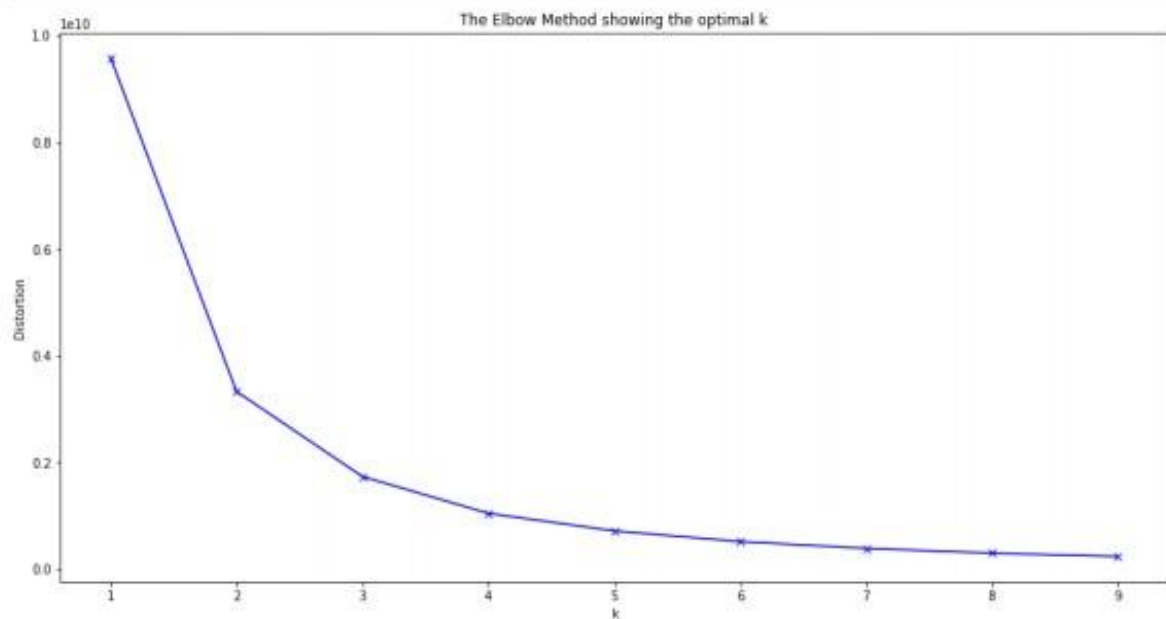
```
In [218]: df.dtypes #All the datatypes are converted into numeric
```

```
Out[218]: QUANTITYORDERED int64
PRICEEACH float64
ORDERLINENUMBER int64
SALES float64
QTR_ID int64
MONTH_ID int64
YEAR_ID int64
MSRP int64
PRODUCTCODE int8
Classic Cars uint8
Motorcycles uint8
Planes uint8
Ships uint8
Trains uint8
Trucks and Buses uint8
Vintage Cars uint8
Large uint8
Medium uint8
Small uint8
dtype: object
```

## Plotting the Elbow Plot to determine the number of clusters.

```
In [219]: distortions = [] # Within Cluster Sum of Squares from the centroid
K = range(1,10)
for k in K:
 kmeanModel = KMeans(n_clusters=k)
 kmeanModel.fit(df)
 distortions.append(kmeanModel.inertia_) #Appending the inertia to the Distortions
```

```
In [220]: plt.figure(figsize=(16,8))
plt.plot(K, distortions, 'bx-')
plt.xlabel('k')
plt.ylabel('Distortion')
plt.title('The Elbow Method showing the optimal k')
plt.show()
```



As the number of k increases Inertia decreases.

Observations: A Elbow can be observed at 3 and after that the curve decreases gradually.

```
In [221]: X_train = df.values #Returns a numpy array.
```

```
In [222]: X_train.shape
```

```
Out[222]: (2823, 19)
```

```
In [223]: model = KMeans(n_clusters=3,random_state=2) #Number of cluster = 3
model = model.fit(X_train) #Fitting the values to create a model.
predictions = model.predict(X_train) #Predicting the cluster values (0,1,or 2)
```

```
In [225]: unique,counts = np.unique(predictions,return_counts=True)
```

```
In [226]: counts = counts.reshape(1,3)
```

```
In [227]: counts_df = pd.DataFrame(counts,columns=['Cluster1','Cluster2','Cluster3'])
```

```
In [228]: counts_df.head()
```

```
Out[228]:
```

|   | Cluster1 | Cluster2 | Cluster3 |
|---|----------|----------|----------|
| 0 | 1083     | 1367     | 373      |

## Visualization

```
In [229]: pca = PCA(n_components=2) #Converting all the features into 2 columns to make it easy to visualize using Principal Component Analysis.
```

## Visualization

```
In [229]: pca = PCA(n_components=2) #Converting all the features into 2 columns to make it easy to visualize using Principal Component Analysis.
```

```
In [230]: reduced_X = pd.DataFrame(pca.fit_transform(X_train),columns=['PCA1','PCA2']) #Creating a DataFrame.
```


```
In [231]: reduced_X.head()
```

Out[231]:

|   | PCA1        | PCA2       |
|---|-------------|------------|
| 0 | -682.488323 | -42.819535 |
| 1 | -787.665502 | -41.694991 |
| 2 | 330.732170  | -26.481208 |
| 3 | 193.040232  | -26.285766 |
| 4 | 1651.532874 | -6.891196  |

```
In [232]: #Plotting the normal Scatter Plot
plt.figure(figsize=(14,10))
plt.scatter(reduced_X['PCA1'],reduced_X['PCA2'])
```

Out[232]: <matplotlib.collections.PathCollection at 0x218dc747880>



```
In [233]: model.cluster_centers_ #Finding the centroids. (3 Centroids in total. Each Array contains a centroids for particular feature)
```

Out[233]: array([[ 3.72031394e+01, 9.52120960e+01, 6.44967682e+00,  
 4.13868425e+03, 2.72022161e+00, 7.09879963e+00,  
 2.00379409e+03, 1.13248384e+02, 5.04469067e+01,  
 3.74884580e-01, 1.15420129e-01, 9.41828255e-02,  
 8.21791320e-02, 1.84672207e-02, 1.16343490e-01,  
 1.98522622e-01, 2.08166817e-17, 1.00000000e+00,  
 -6.66133815e-16],  
 [ 3.08302853e+01, 7.00755230e+01, 6.67300658e+00,  
 2.12409474e+03, 2.71762985e+00, 7.09509876e+00,  
 2.00381127e+03, 7.84784199e+01, 6.24871982e+01,  
 2.64813460e-01, 1.21433797e-01, 1.29480614e-01,  
 1.00219459e-01, 3.87710315e-02, 9.21726408e-02,  
 2.53108990e-01, 6.93889390e-18, 6.21799561e-02,  
 9.37820044e-01],  
 [ 4.45871314e+01, 9.98931099e+01, 5.75603217e+00,  
 7.09596863e+03, 2.71845576e+00, 7.06434316e+00,  
 2.00380900e+03, 1.45823056e+02, 3.14959786e+01,  
 5.33512064e-01, 1.07238606e-01, 7.23860590e-02,  
 2.14477212e-02, 1.07238606e-02, 1.31367292e-01,  
 1.23324397e-01, 4.20911528e-01, 5.79088472e-01,  
 5.55111512e-17]])



```

9.37820044e-01],
[4.45871314e+01, 9.98931099e+01, 5.75603217e+00,
 7.09596863e+03, 2.71045576e+00, 7.06434316e+00,
 2.00389008e+03, 1.45823056e+02, 3.14959786e+01,
 5.33512064e-01, 1.07238606e-01, 7.23860590e-02,
 2.14477212e-02, 1.07238606e-02, 1.31367292e-01,
 1.23324397e-01, 4.20911528e-01, 5.79088472e-01,
 5.55111512e-17]])

```


```
In [234]: reduced_centers = pca.transform(model.cluster_centers_) #Transforming the centroids into 3 in x and y coordinates
```

```
In [235]: reduced_centers
```

```
Out[235]: array([[5.84994044e+02, -4.36786931e+00],
 [-1.43005891e+03, 2.60041009e+00],
 [3.54247180e+03, 3.15185487e+00]])
```

```
In [236]: plt.figure(figsize=(14,10))
plt.scatter(reduced_X['PCA1'],reduced_X['PCA2'])
plt.scatter(reduced_centers[:,0],reduced_centers[:,1],color='black',marker='x',s=300) #Plotting the centroids
```

```
Out[236]: <matplotlib.collections.PathCollection at 0x218deb6e220>
```



```
In [237]: reduced_X['Clusters'] = predictions #Adding the Clusters to the reduced dataframe.
```

```
In [238]: reduced_X.head()
```

```
Out[238]:
```

|   | PCA1        | PCA2       | Clusters |
|---|-------------|------------|----------|
| 0 | -682.488323 | -42.819535 | 1        |
| 1 | -787.665502 | -41.694991 | 1        |
| 2 | 330.732170  | -26.481208 | 0        |
| 3 | 193.040232  | -26.285766 | 0        |
| 4 | 1651.532874 | -6.891196  | 0        |

```
In [239]: #Plotting the clusters
plt.figure(figsize=(14,10))
taking the cluster number and first column taking the same cluster number and second
column Assigning the color
plt.scatter(reduced_X[reduced_X['Clusters'] == 0].loc[:, 'PCA1'],reduced_X[reduced_X['Clusters'] == 0].loc[:, 'PCA2'],
color='slateblue')
plt.scatter(reduced_X[reduced_X['Clusters'] == 1].loc[:, 'PCA1'],reduced_X[reduced_X['Clusters'] == 1].loc[:, 'PCA2'],
color='springgreen')
plt.scatter(reduced_X[reduced_X['Clusters'] == 2].loc[:, 'PCA1'],reduced_X[reduced_X['Clusters'] == 2].loc[:, 'PCA2'],
color='indigo')

plt.scatter(reduced_centers[:,0],reduced_centers[:,1],color='black',marker='x',s=300)
```

```
In [237]: reduced_X['Clusters'] = predictions #Adding the Clusters to the reduced dataframe.
```

```
In [238]: reduced_X.head()
```

```
Out[238]:
```

|   | PCA1        | PCA2       | Clusters |
|---|-------------|------------|----------|
| 0 | -682.488323 | -42.819535 | 1        |
| 1 | -787.665502 | -41.694991 | 1        |
| 2 | 330.732170  | -26.481208 | 0        |
| 3 | 193.040232  | -26.285766 | 0        |
| 4 | 1651.532874 | -6.891196  | 0        |

```
In [239]: #Plotting the clusters
plt.figure(figsize=(14,10))
taking the cluster number and first column taking the same cluster number and second
column Assigning the color
plt.scatter(reduced_X[reduced_X['Clusters'] == 0].loc[:, 'PCA1'], reduced_X[reduced_X['Clusters'] == 0].loc[:, 'PCA2'],
color='slateblue')
plt.scatter(reduced_X[reduced_X['Clusters'] == 1].loc[:, 'PCA1'], reduced_X[reduced_X['Clusters'] == 1].loc[:, 'PCA2'],
color='springgreen')
plt.scatter(reduced_X[reduced_X['Clusters'] == 2].loc[:, 'PCA1'], reduced_X[reduced_X['Clusters'] == 2].loc[:, 'PCA2'],
color='indigo')

plt.scatter(reduced_centers[:,0], reduced_centers[:,1], color='black', marker='x', s=300)
```

```
Out[239]: <matplotlib.collections.PathCollection at 0x218dce9e1f0>
```

